Identity Resolver - Historical Issues List
All Resolved – definitely emerged over time…

Incident	Resolution	Humorous Image
Same Batch Jobs being processed on different webservers, getting instances where a clients file is being picked up and processed multiple times! This leading to partial processing and a confused FTP server! Critical entry point for new clients before they integrate with the API. Fix ASAP!	It appears that we cannot trust the threading implementation in Perl for our batch file processing. There are (eventually) less clients on batch, so lets limit to one instance rather than three, rather than go through all the pain of developing a thread safe solution…	
Clients have reported that old search results are being served to them, or results they didn't ask for! This is happening intermittently, can we please investigate our page and query caching?	Our caching implementation is supposed to be smart enough to expire pages and query results for the whole cluster, but it isn't. Given the limited number of website clients, lets give each Webserver has its own cache, rather than a distributed one.	
We have just noted that if an attacker managed to penetrate our external firewall in front of the load balancer, our lack of internal firewall to other systems would leave us exposed to attack. Also, we can track internal traffic much more effectively.	Lets install an internal firewall our to Credit Search and Bank Account Check services. This should also be load balanced to protect other systems from surges in demand for ID Resolver…	

Identity Resolver - Historical Issues List
All Resolved – definitely emerged over time…

Each customer has results they wish to retrieve historically and application configuration specific to their needs. There was an incident on live the other day when an accidental deletion lead to a clients config being rebuilt…	It would seem sensible to add customer results and config into a failover database server, as well as the audit and billing data. Then at least we can request a restore of configuration and results, or at least look at it, without going to get it from physical media…	
Now we have added individual caches per webserver, unfortunately clients are now receiving different results for cached responses. This is resulting in inconsistent results for credit scoring.	It comes at a cost but we can enable sticky sessions on the load balancer, therefore a client will stick to one webserver, and you can deal with their traffic in an isolated way, and redirect everyone else to the other two…	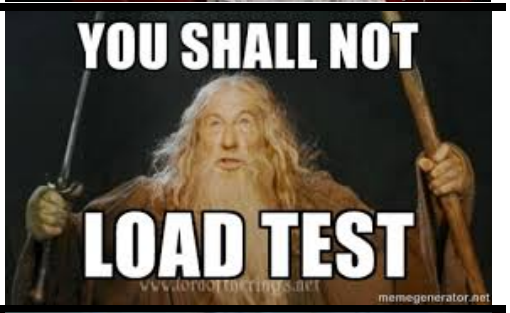
There have been a number of gaps within the billing data generated by ID Resolver, when reconciled against its interactions with internal and external data providers. Can we investigate the effectiveness of replication and have somewhere to compare the support database too?	Lets replicate the billing and audit data from the live instances to the failover instance, then onto the support instance and apply some monitoring on the live to failover replication.